
WHY AND HOW TO LEVERAGE THE POWER AND SIMPLICITY OF SQL ON APACHE FLINK®

FABIAN HUESKE, SOFTWARE ENGINEER

ABOUT ME

- Apache Flink PMC member & ASF member
 - Contributing since day 1 at TU Berlin
 - Focusing on Flink's relational APIs since ~2 years
- Co-author of “Stream Processing with Apache Flink”
 - Work in progress...
- Co-founder & Software Engineer at data Artisans



ABOUT DATA ARTISANS



Original creators of
Apache Flink®

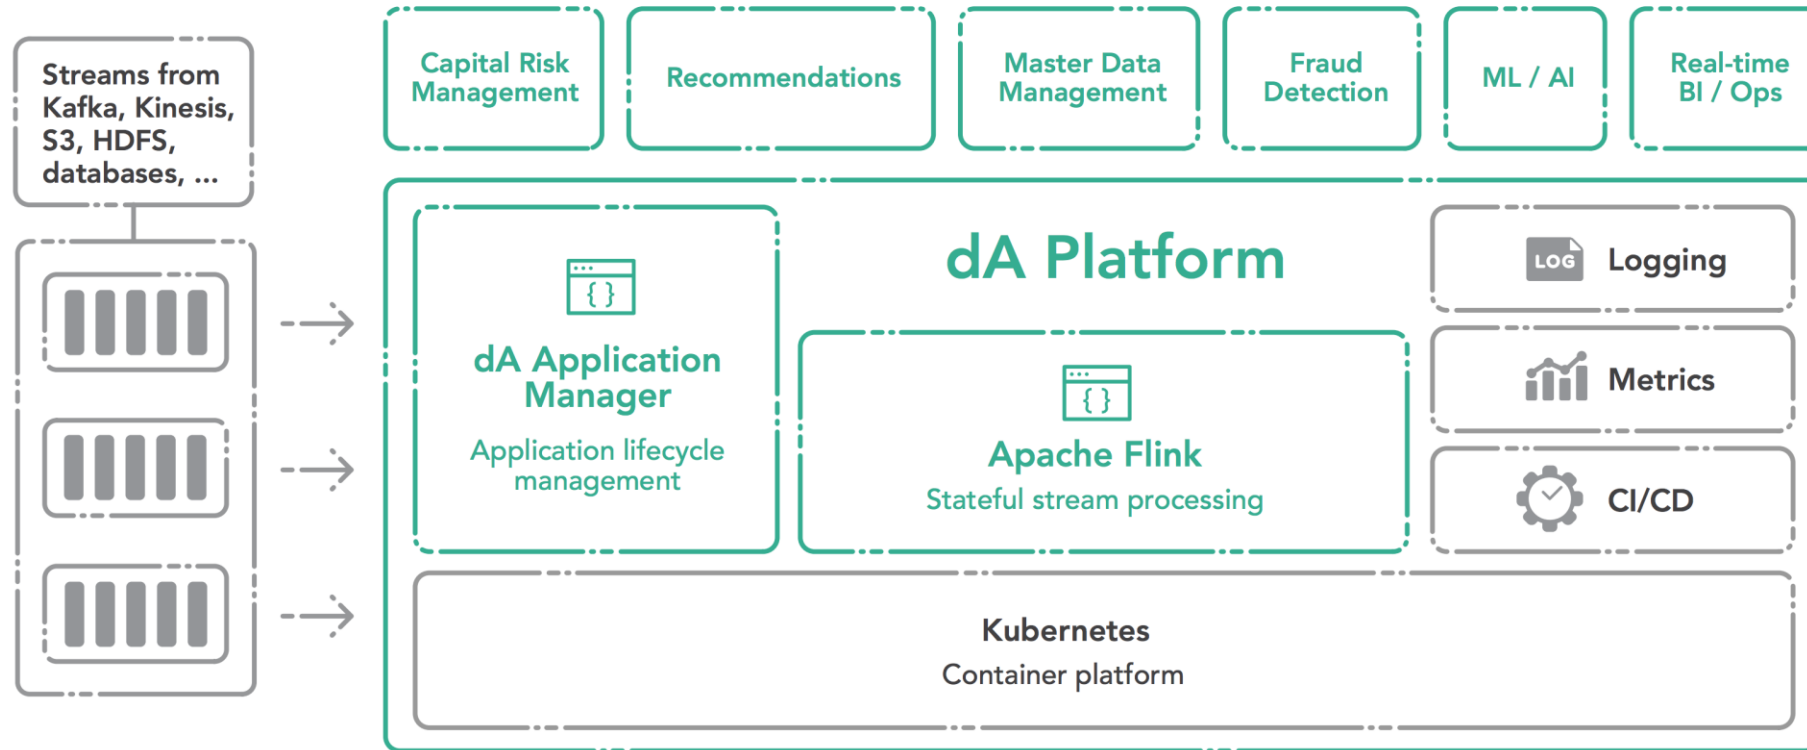


PLATFORM

Open Source Apache Flink
+ dA Application Manager



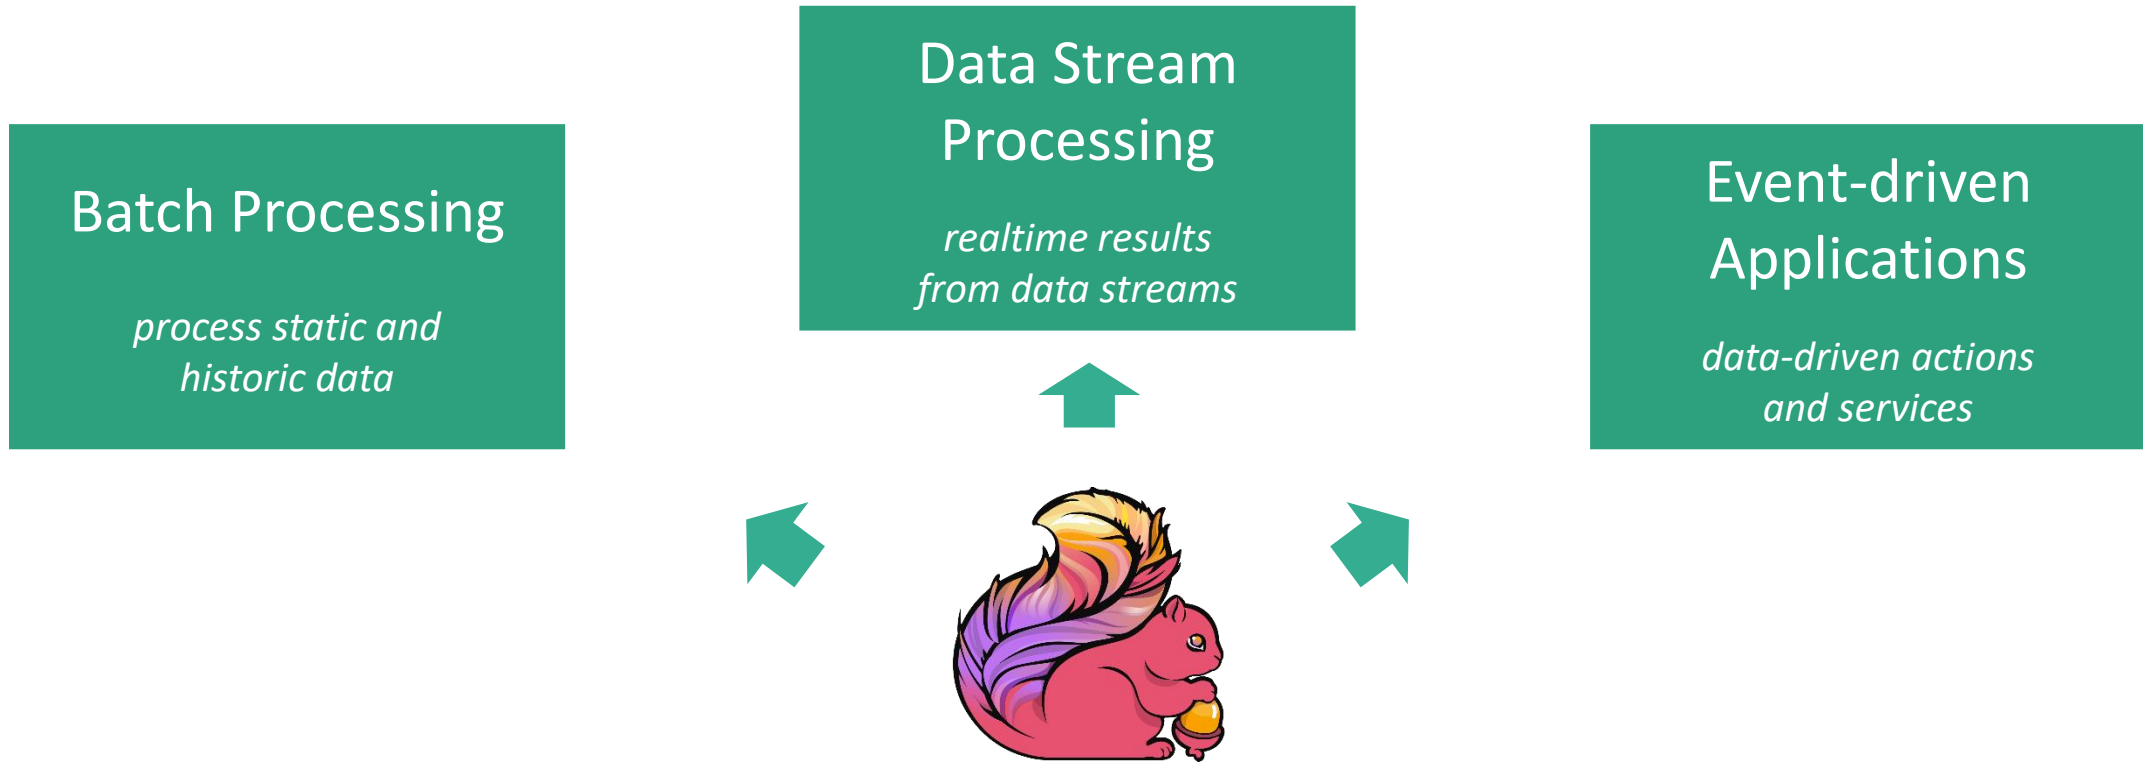
DA PLATFORM



data-artisans.com/download



WHAT IS APACHE FLINK?

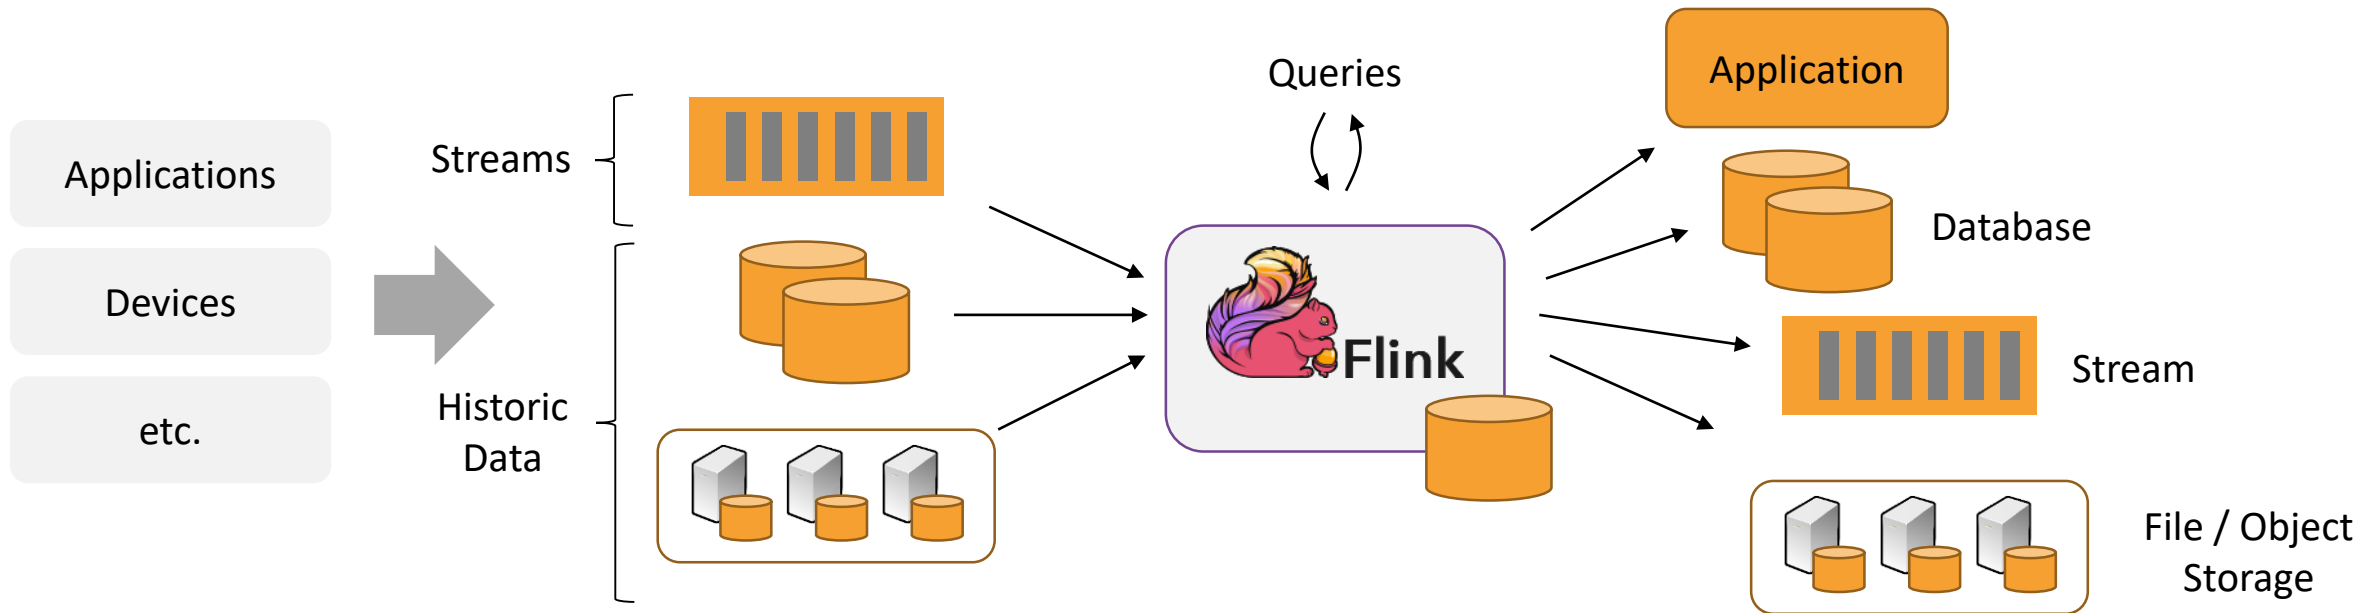


Stateful Computations Over Data Streams



WHAT IS APACHE FLINK?

Stateful computations over streams
real-time and historic
fast, scalable, fault tolerant, in-memory,
event time, large state, exactly-once



POWERED BY APACHE FLINK



FLINK'S POWERFUL ABSTRACTIONS

Layered abstractions to
navigate simple to complex use cases

High-level
Analytics API

SQL / Table API (dynamic tables)

```
SELECT room, TUMBLE_END(rowtime, INTERVAL '1' HOUR), AVG(temp)
FROM sensors
GROUP BY TUMBLE(rowtime, INTERVAL '1' HOUR), room
```

Stream- & Batch
Data Processing

DataStream API (streams, windows)

```
val stats = stream
  .keyBy("sensor")
  .timeWindow(Time.seconds(5))
  .sum((a, b) -> a.add(b))
```

Stateful Event-
Driven Applications

Process Function (events, state, time)

```
def processElement(event: MyEvent, ctx: Context, out: Collector[Result]) = {
  // work with event and state
  (event, state.value) match { ... }

  out.collect(...) // emit events
  state.update(...) // modify state

  // schedule a timer callback
  ctx.timerService.registerEventTimeTimer(event.timestamp + 500)
}
```



APACHE FLINK'S RELATIONAL APIS

ANSI SQL

```
SELECT user, COUNT(url) AS cnt  
FROM clicks  
GROUP BY user
```

LINQ-style Table API

```
tableEnvironment  
    .scan("clicks")  
    .groupBy('user')  
    .select('user', 'url.count as 'cnt')
```

Unified APIs for batch & streaming data

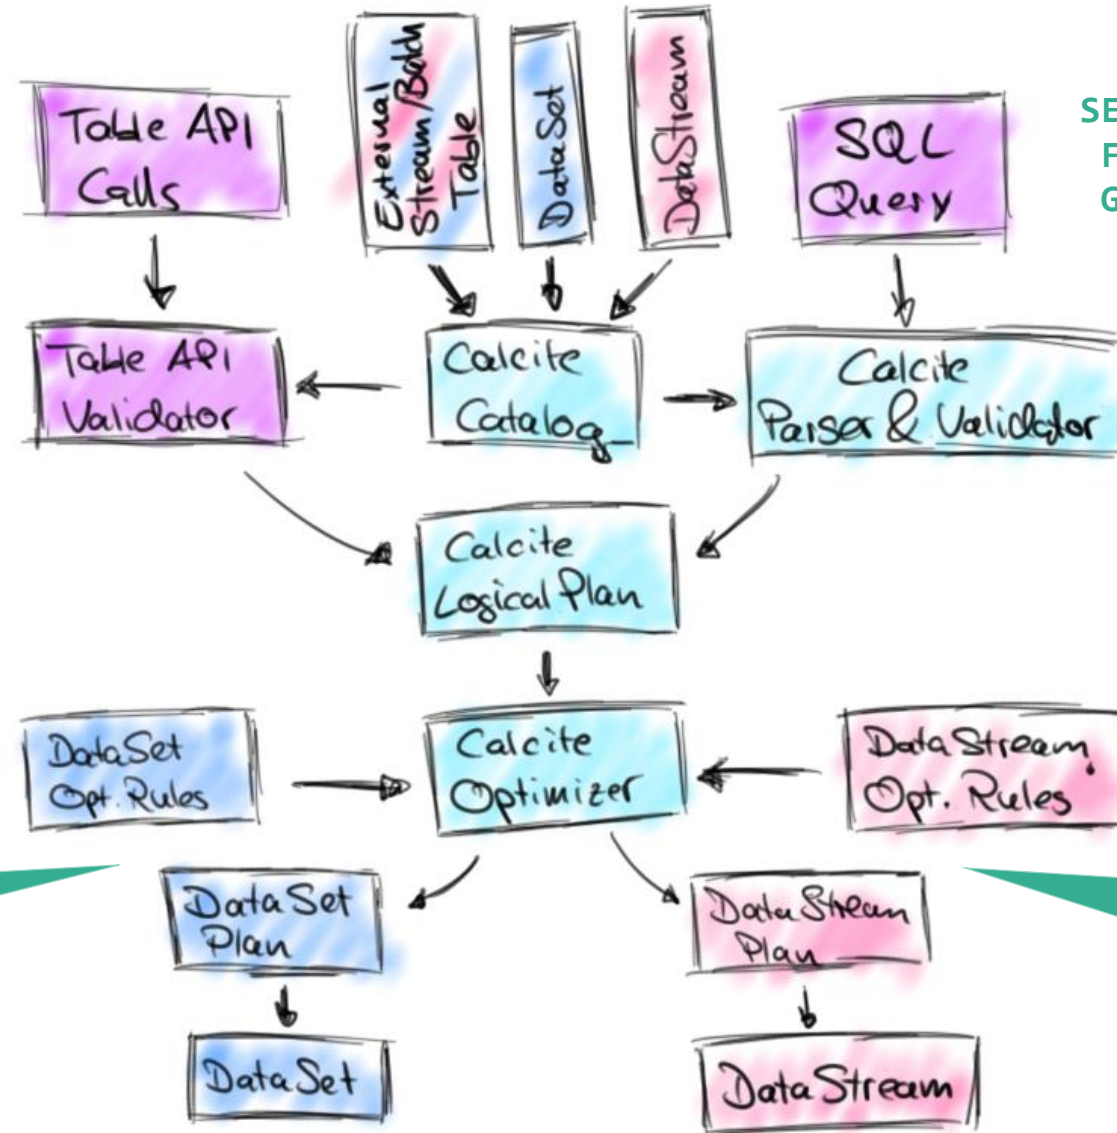
*A query specifies exactly the same result
regardless whether its input is
static batch data or streaming data.*



QUERY TRANSLATION

```
tableEnvironment  
  .scan("clicks")  
  .groupBy('user')  
  .select('user', 'url.count as 'cnt')
```

```
SELECT user, COUNT(url) AS cnt  
FROM clicks  
GROUP BY user
```

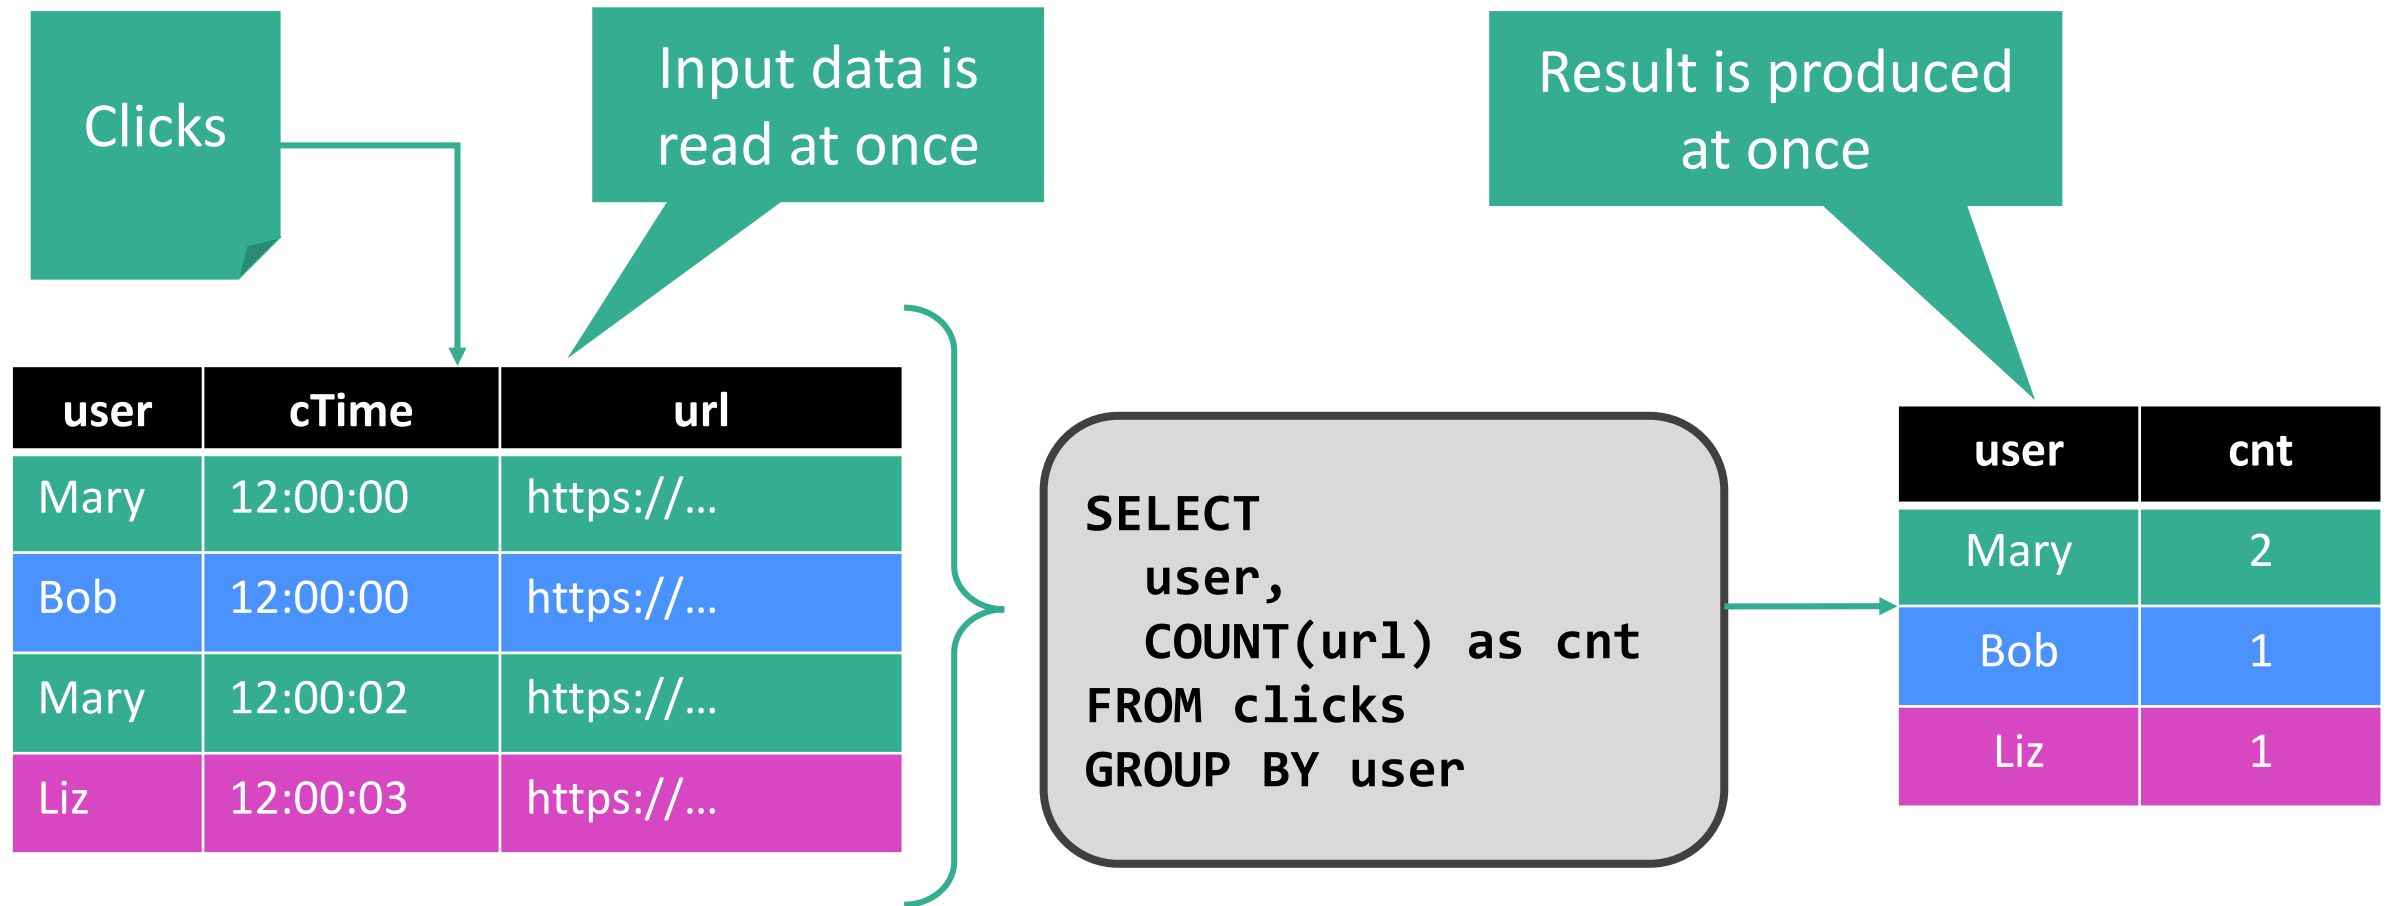


Input data is
bounded
(batch)

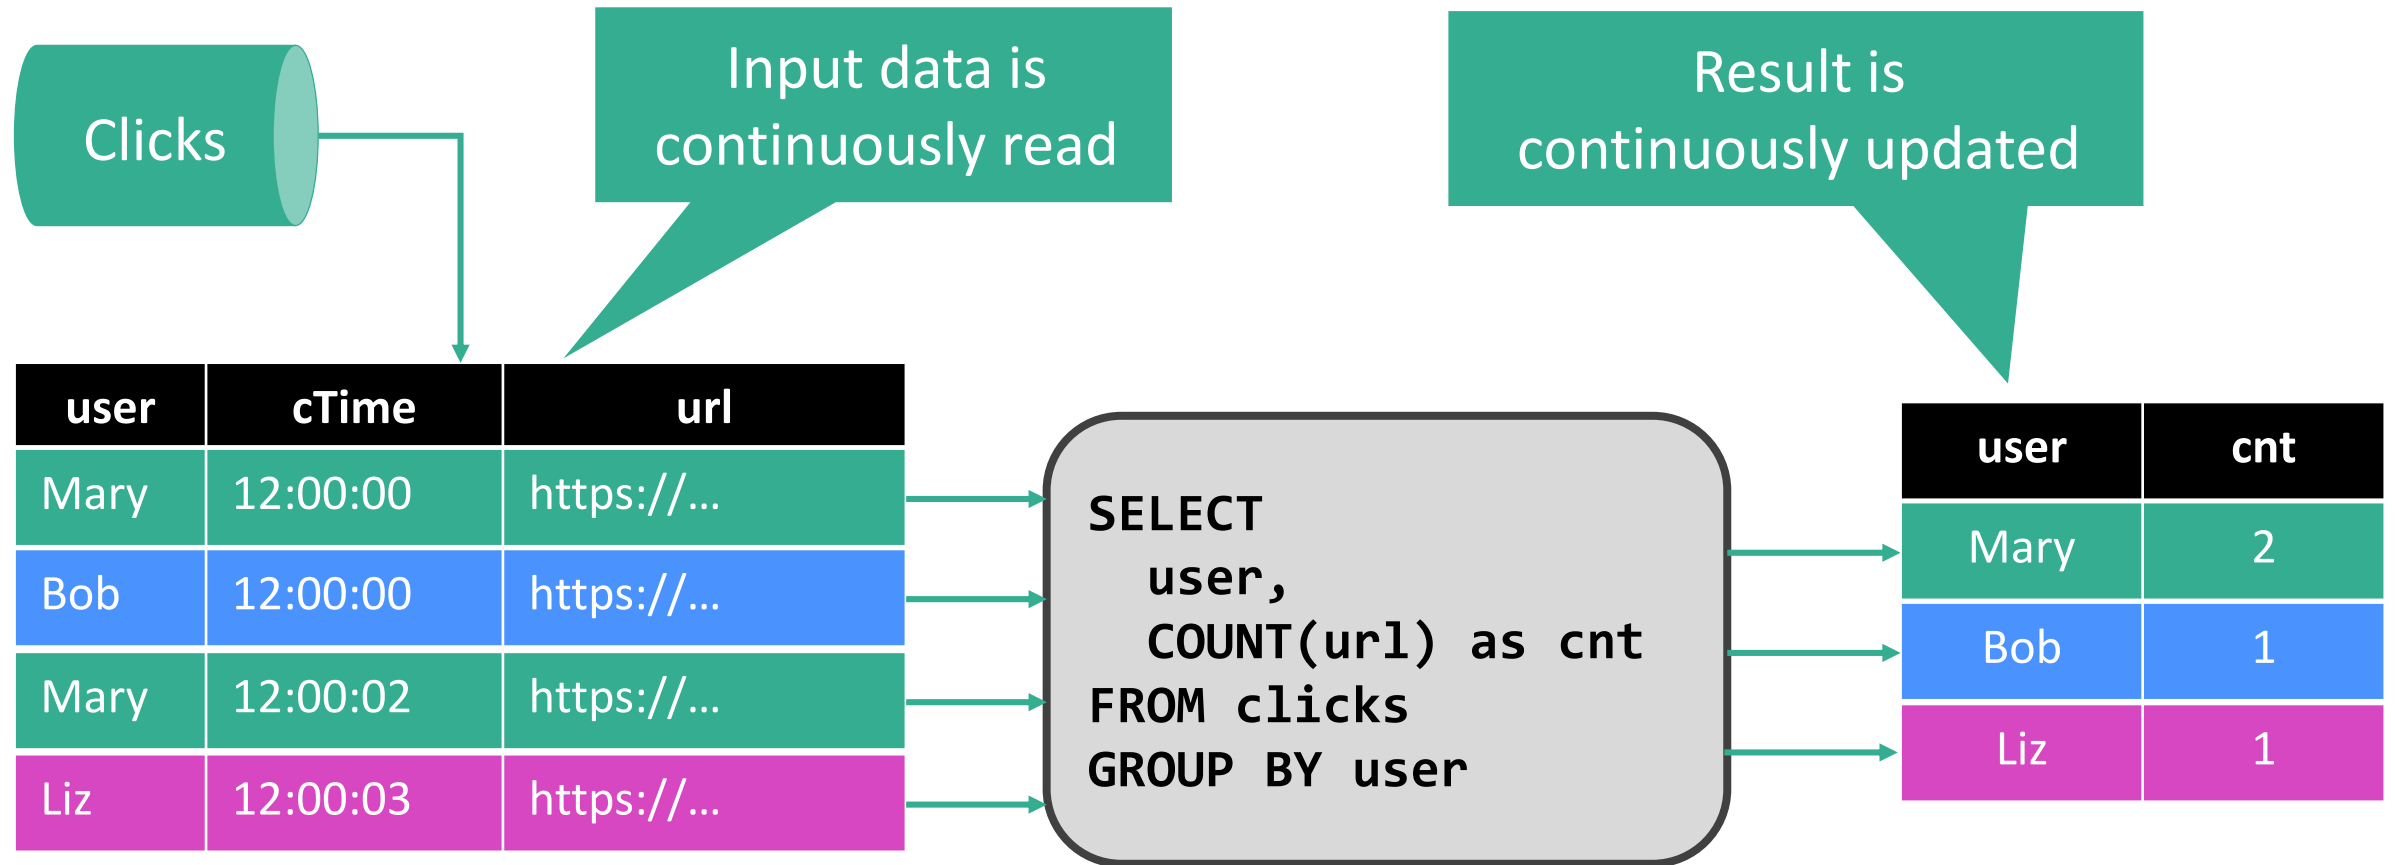
Input data is
unbounded
(streaming)



WHAT IF “CLICKS” IS A FILE?



WHAT IF “CLICKS” IS A STREAM?




The result is the same!



WHY IS STREAM-BATCH UNIFICATION IMPORTANT?

- Usability
 - ANSI SQL syntax: No custom “StreamSQL” syntax.
 - ANSI SQL semantics: No stream-specific results.
- Portability
 - Run the same query on *bounded* and *unbounded* data
 - Run the same query on *recorded* and *real-time* data



- How can we achieve SQL semantics on streams?



DATABASE SYSTEMS RUN QUERIES ON STREAMS

- Materialized views (MV) are similar to regular views, but persisted to disk or memory
 - Used to speed-up analytical queries
 - MVs need to be updated when the base tables change
- MV maintenance is very similar to SQL on streams
 - Base table updates are a stream of DML statements
 - MV definition query is evaluated on that stream
 - MV is query result and continuously updated



CONTINUOUS QUERIES IN FLINK

- Core concept is a “*Dynamic Table*”
 - Dynamic tables are changing over time
- Queries on dynamic tables
 - produce new dynamic tables (which are updated based on input)
 - do not terminate
- Stream ↔ Dynamic table conversions



STREAM ↔ DYNAMIC TABLE CONVERSIONS

- Append Conversions
 - Records are only inserted/appended
- Upsert Conversions
 - Records are inserted/updated/deleted
 - Records have a (composite) unique key
- Changelog Conversions
 - Records are inserted/updated/deleted



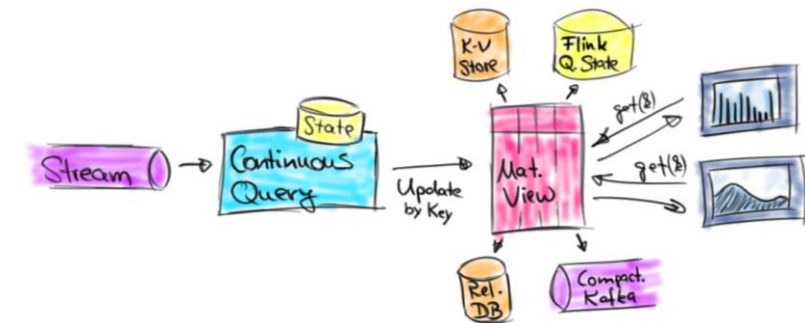

SQL FEATURE SET IN FLINK 1.5.0

- SELECT FROM WHERE
- GROUP BY / HAVING
 - Non-windowed, TUMBLE, HOP, SESSION windows
- JOIN
 - Windowed INNER, LEFT / RIGHT / FULL OUTER JOIN
 - Non-windowed INNER JOIN
- Scalar, aggregation, table-valued UDFs
- SQL CLI Client (beta)
- [streaming only] OVER / WINDOW
 - UNBOUNDED / BOUNDED PRECEDING
- [batch only] UNION / INTERSECT / EXCEPT / IN / ORDER BY



WHAT CAN I BUILD WITH THIS?

- Data Pipelines
 - Transform, aggregate, and move events in real-time
- Low-latency ETL
 - Convert and write streams to file systems, DBMS, K-V stores, indexes, ...
 - Ingest appearing files to produce streams
- Stream & Batch Analytics
 - Run analytical queries over bounded and unbounded data
 - Query and compare historic and real-time data
- Power Live Dashboards
 - Compute and update data to visualize in real-time



THE NEW YORK TAXI RIDES DATA SET

- The New York City Taxi & Limousine Commission provides a public data set about past taxi rides in New York City
- We can derive a streaming table from the data
- Table: **TaxiRides**

<code>rideId:</code>	<code>BIGINT</code>	<code>// ID of the taxi ride</code>
<code>isStart:</code>	<code>BOOLEAN</code>	<code>// flag for pick-up (true) or drop-off (false) event</code>
<code>lon:</code>	<code>DOUBLE</code>	<code>// longitude of pick-up or drop-off location</code>
<code>lat:</code>	<code>DOUBLE</code>	<code>// latitude of pick-up or drop-off location</code>
<code>rowtime:</code>	<code>TIMESTAMP</code>	<code>// time of pick-up or drop-off event</code>



IDENTIFY POPULAR PICK-UP / DROP-OFF LOCATIONS

- Compute *every 5 minutes* for *each location* the *number of departing and arriving taxis* of *the last 15 minutes*.

```
SELECT cell,  
       isStart,  
       HOP_END(rowtime, INTERVAL '5' MINUTE, INTERVAL '15' MINUTE) AS hopEnd,  
       COUNT(*) AS cnt  
FROM (SELECT rowtime, isStart, toCellId(lon, lat) AS cell  
      FROM TaxiRides)  
GROUP BY cell,  
         isStart,  
         HOP(rowtime, INTERVAL '5' MINUTE, INTERVAL '15' MINUTE)
```



AVERAGE RIDE DURATION PER PICK-UP LOCATION



- *Join start ride and end ride events on rideId and compute average ride duration per pick-up location.*

```
SELECT pickUpCell,  
       AVG(TIMESTAMPDIFF(MINUTE, e.rowtime, s.rowtime)) AS avgDuration  
FROM (SELECT rideId, rowtime, toCellId(lon, lat) AS pickUpCell  
      FROM TaxiRides  
      WHERE isStart) s  
JOIN  
      (SELECT rideId, rowtime  
      FROM TaxiRides  
      WHERE NOT isStart) e  
ON s.rideId = e.rideId AND  
   e.rowtime BETWEEN s.rowtime AND s.rowtime + INTERVAL '1' HOUR  
GROUP BY pickUpCell
```



BUILDING A DASHBOARD

```
SELECT cell,  
       isStart,  
       HOP_END(rowtime, INTERVAL '5' MINUTE, INTERVAL '15' MINUTE) AS hopEnd,  
       COUNT(*) AS cnt  
FROM (SELECT rowtime, isStart, toCellId(lon, lat) AS cell  
      FROM TaxiRides)  
GROUP BY cell,  
         isStart,  
         HOP(rowtime, INTERVAL '5' MINUTE, INTERVAL '15' MINUTE)
```



SOUNDS GREAT! HOW CAN I USE IT?

- ATM, SQL queries must be embedded in Java/Scala code ☹️
 - Tight integration with DataStream and DataSet APIs
- Community focused on internals (until Flink 1.4.0)
 - Operators, types, built-in functions, extensibility (UDFs, extern. catalog)
 - Proven at scale by Alibaba, Huawei, and Uber
 - All built their own submission system & connectors library
- Community neglected user interfaces
 - No query submission client, no CLI
 - No integration with common catalog services
 - Limited set of TableSources and TableSinks



COMING IN FLINK 1.5.0 - SQL CLI

Demo Time!

That's a nice toy, but ...

... can I use it for anything serious?



FLIP-24 – A SQL QUERY SERVICE

- REST service to submit & manage SQL queries
 - SELECT ...
 - INSERT INTO SELECT ...
 - CREATE MATERIALIZED VIEW ...
- Serve results of “SELECT ...” queries
- Provide a table catalog (integrated with external catalogs)
- Use cases
 - Data exploration with notebooks like Apache Zeppelin
 - Access to real-time data from applications
 - Easy data routing / ETL from management consoles



CHALLENGE: SERVE DYNAMIC TABLES

Unbounded input yields unbounded results

(Serving bounded results is easy)

```
SELECT user, url
FROM clicks
WHERE url LIKE '%xyz.com'
```

```
SELECT user, COUNT(url) AS cnt
FROM clicks
GROUP BY user
```

Append-only Table

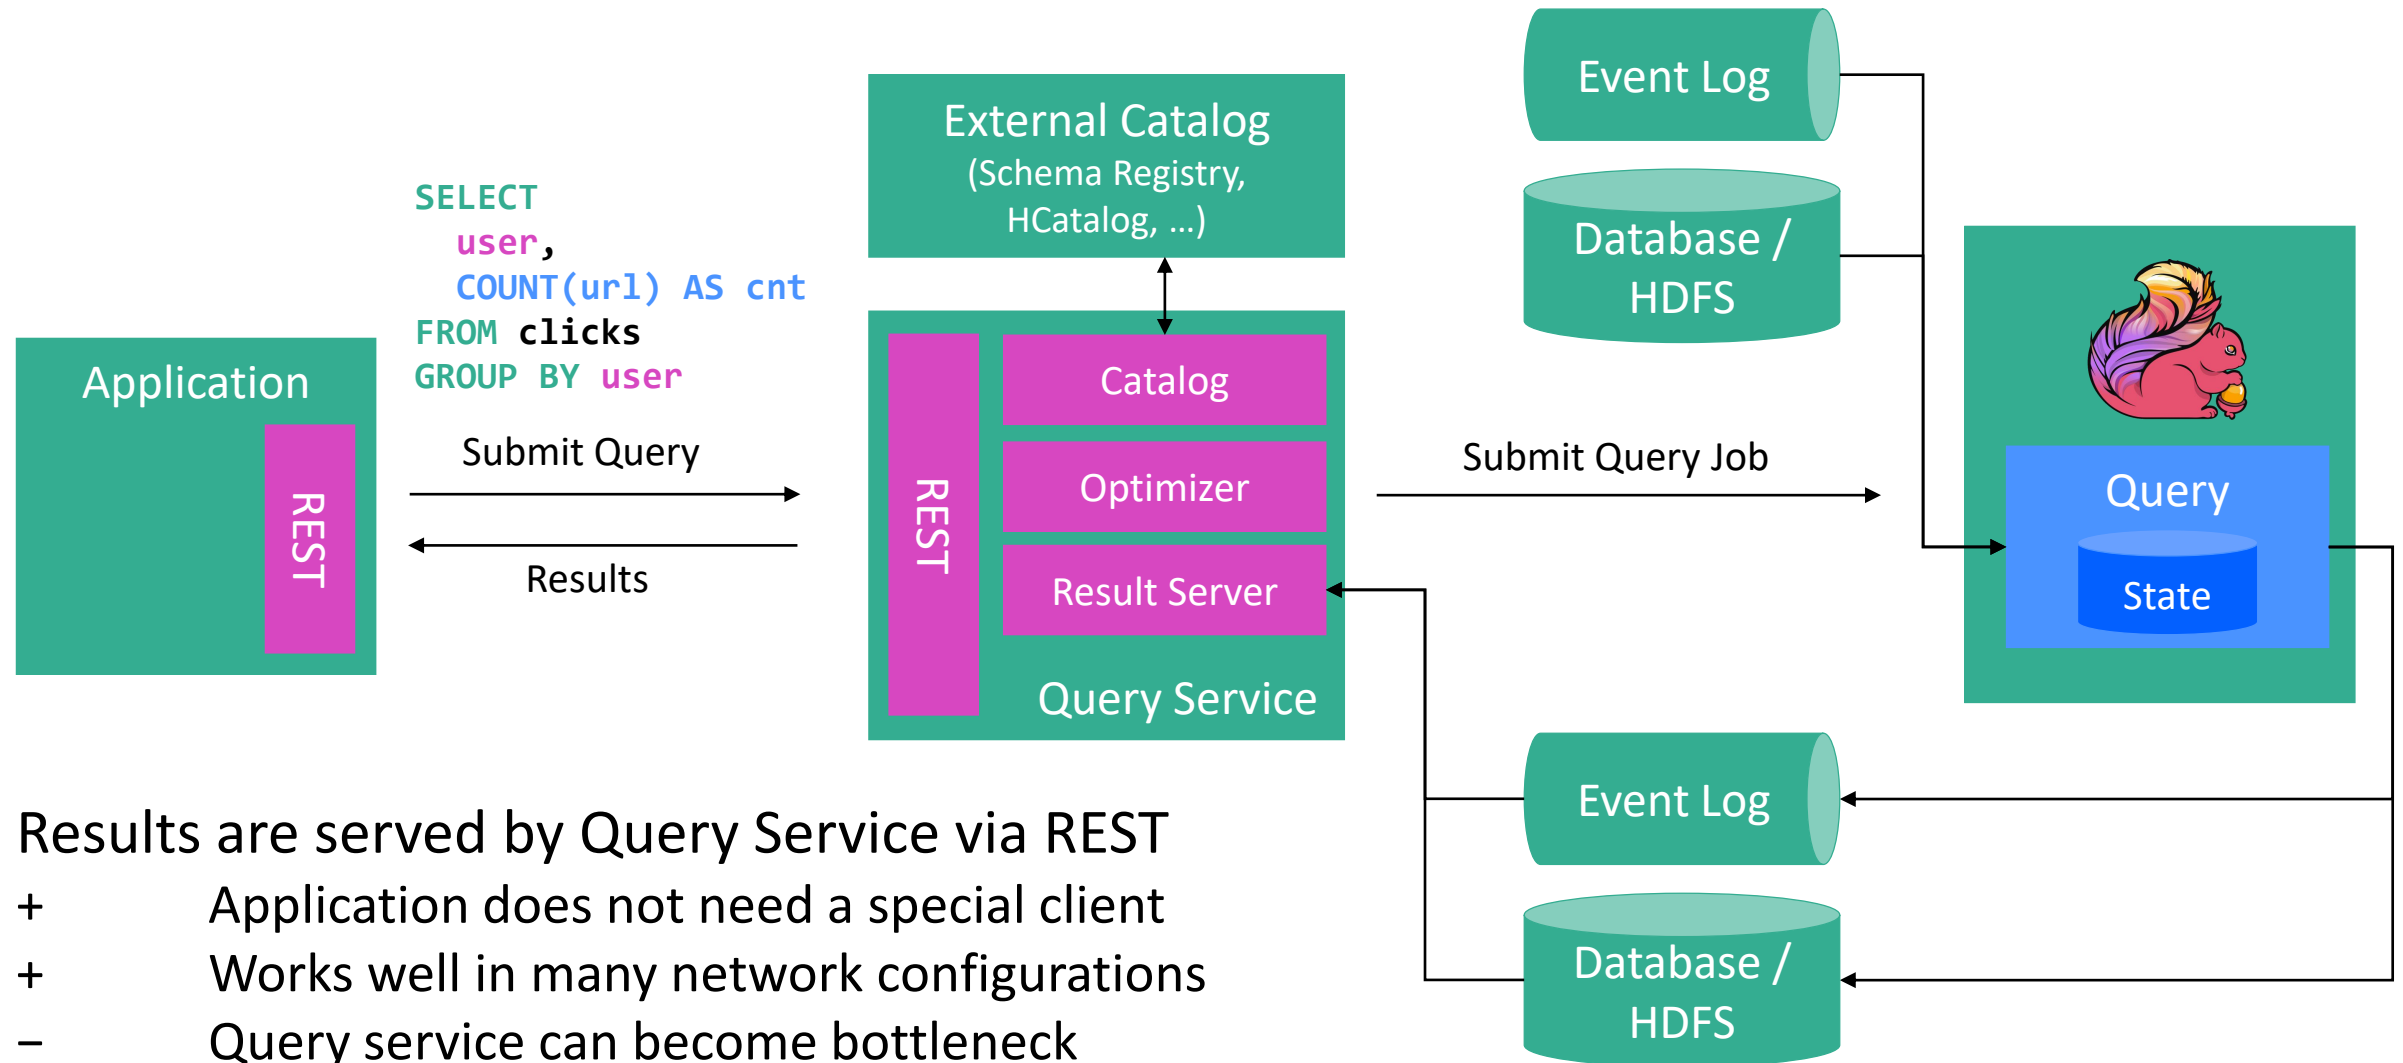
- Result rows are never changed
- Consume, buffer, or drop rows

Continuously updating Table

- Result rows can be updated or deleted
- Consume changelog or periodically query result table
- Result table must be maintained somewhere



FLIP-24 – A SQL QUERY SERVICE

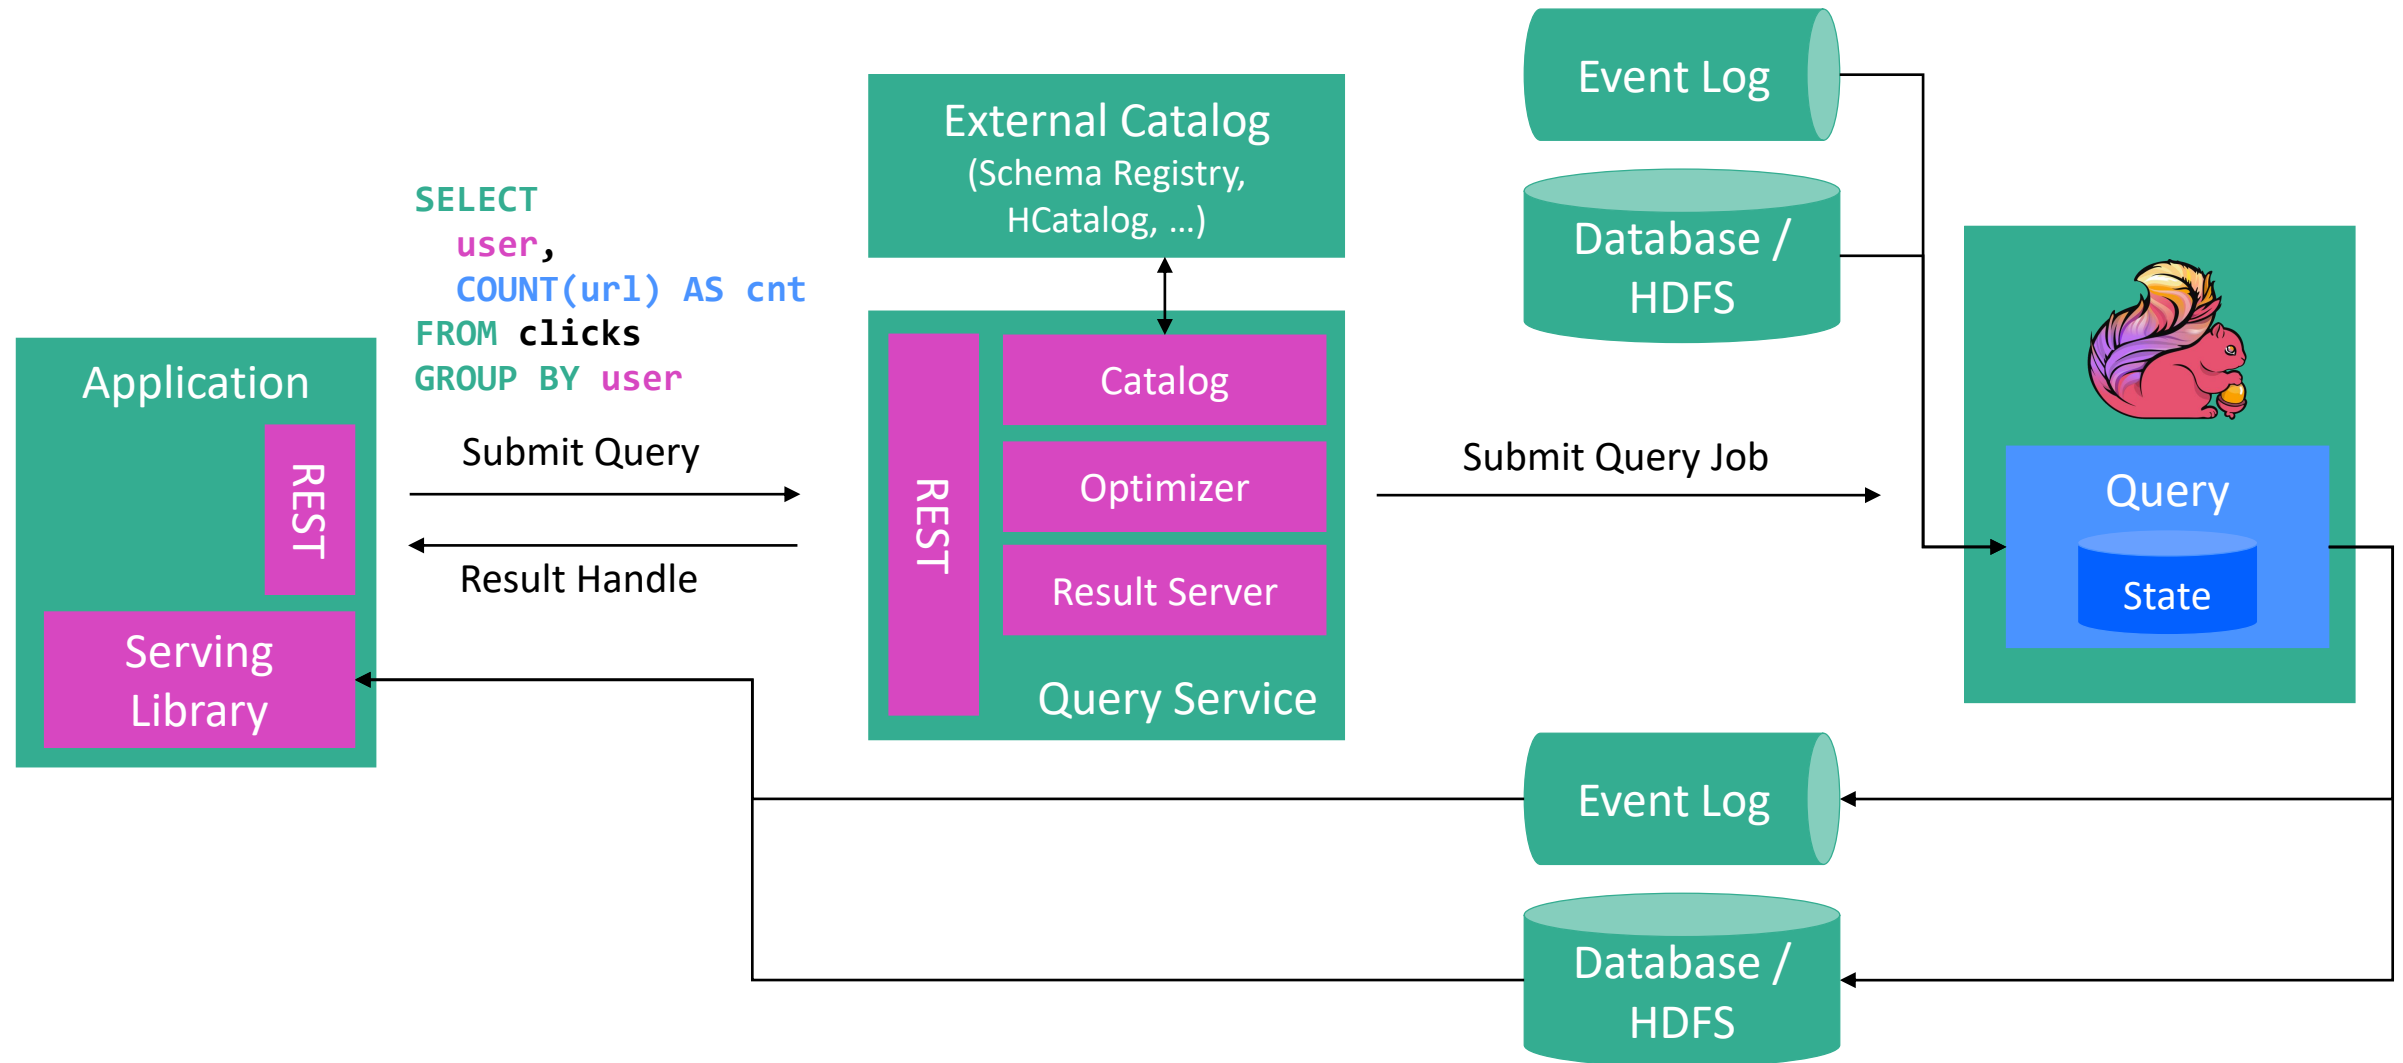


Results are served by Query Service via REST

- + Application does not need a special client
- + Works well in many network configurations
- Query service can become bottleneck



FLIP-24 – A SQL QUERY SERVICE



WE WANT YOUR FEEDBACK!

- The design of SQL Query Service is not final yet.
- Check out FLIP-24 and FLINK-7594
- Share your ideas and feedback and discuss on JIRA or dev@flink.apache.org.




SUMMARY

- Unification of stream and batch is important.
- Flink's SQL solves many streaming and batch use cases.
- Runs in production at Alibaba, Uber, and others.
- The community is working on improving user interfaces.
- Get involved, discuss, and contribute!



THANK YOU!



Available on O'Reilly Early Release!

dataArtisans

THANK YOU!

@fhueske

@dataArtisans

@ApacheFlink

WE ARE HIRING

data-artisans.com/careers

dataArtisans